

Word Completion

Thomas Graf

Stony Brook University
lin120@thomasgraf.net

LIN 220, Spring 2019
Lecture 2

Word completion: a simple problem(?)

- ▶ When you write something on your phone, it automatically suggests words while you're typing them.
- ▶ Seems easy, but it's **fairly intricate**.

Key Insights/Skills

- ▶ transition probabilities with n-grams
- ▶ “More than 1, but less than 5.”
- ▶ efficient data structures: tries/prefix trees

Attempt 1: simple lookup

- ▶ To make suggestions for completions, we only need an English/German/... dictionary.

```
1  # load English dictionary from nltk package
2  from nltk.corpus import words
3  # and see if "test" is in the list
4  "test" in words.words()
5  >>> True
```

- ▶ Given word **w**, the set of possible completions for **w** consists of all listed words that start with **w**:

$$\text{completions}(\mathbf{w}) = \{\mathbf{w}' \mid \mathbf{w}' \text{ is a word of English, and } \mathbf{w}' \text{ starts with } \mathbf{w}\}$$

Attempt 1 [cont.]

- completions(**w**) is easily computed in Python.

```
1  def completions(word, wordlist):  
2      """Return set of all known completions for word."""  
3      return {comp for comp in wordlist  
4              if comp.startswith(word)}
```

```
1  completions("testing", words.words())  
2  >>> {"testing", "testingly"}
```

Evaluating attempt 1

- Is this a **good solution?**

Pro	Contra

```
1 completions("test", words.words())
2 >>> {"test", "testa", "testable", "testacean", ...}
```

Evaluating attempt 1

- Is this a **good solution?**

Pro	Contra
simple	

```
1 completions("test", words.words())
2 >>> {"test", "testa", "testable", "testacean", ...}
```

Evaluating attempt 1

- Is this a **good solution?**

Pro	Contra
simple	too many completions

```
1 completions("test", words.words())
2 >>> {"test", "testa", "testable", "testacean", ...}
```

Evaluating attempt 1

- Is this a **good solution?**

Pro	Contra
simple	too many completions unlikely completions

```
1 completions("test", words.words())
2 >>> {"test", "testa", "testable", "testacean", ...}
```


Improving attempt 1

- We can limit the number of suggestions, if we use a **list** instead of a **set**.

```
1  def completions(word, wordlist):  
2      """Return list of all known completions for word."""  
3      return [comp for comp in wordlist  
4              if comp.startswith(word)]  
5  
6  completions("test", words.words())[:3]  
7  >>> ["test", "testa", "testable"]
```

- But this still includes unlikely completions like *testa*.
- We need **probabilities**!

Probability = frequency (?)

- ▶ We only want to show the **most likely completions**.
- ▶ But what is most likely?
- ▶ **Naive proposal:** the most frequent word!

Probability = frequency (?)

- ▶ We only want to show the **most likely completions**.
- ▶ But what is most likely?
- ▶ **Naive proposal:** the most frequent word!

Um, but how do we figure out what is most common?

How to find common words

- ▶ **Task:** determine which words are common
- ▶ **Solution:**
 - 1 Collect sufficiently large sample of texts
 - 2 For each word (**type**), count how often it occurs in the entire sample (= its number of **tokens**).
 - 3 Calculate the **frequency** of the word in the sample:

$$\text{freq}(\text{word}, \text{sample}) = \frac{\text{number of tokens of word}}{\text{word length of whole sample}}$$

Types vs tokens

We have to distinguish **word types** (*a, the, Mary, red, ...*) from their **word tokens**, which are the instances of a specific word type. For instance, the type “word” has 4 tokens in this box.

How to find common words

- ▶ **Task:** determine which words are common
- ▶ **Solution:**
 - 1 Collect sufficiently large sample of texts
 - 2 For each word (**type**), count how often it occurs in the entire sample (= its number of **tokens**).
 - 3 Calculate the **frequency** of the word in the sample:

$$\text{freq}(\text{word}, \text{sample}) = \frac{\text{number of tokens of word}}{\text{word length of whole sample}}$$

Types vs tokens

We have to distinguish **word types** (*a*, *the*, *Mary*, *red*, ...) from their **word tokens**, which are the instances of a specific word type. For instance, the type “word” has 4 tokens in this box.

How to find common words

- ▶ **Task:** determine which words are common
- ▶ **Solution:**
 - 1 Collect sufficiently large sample of texts
 - 2 For each word (**type**), count how often it occurs in the entire sample (= its number of **tokens**).
 - 3 Calculate the **frequency** of the word in the sample:

$$\text{freq}(\text{word}, \text{sample}) = \frac{\text{number of tokens of word}}{\text{word length of whole sample}}$$

Types vs tokens

We have to distinguish **word types** (*a, the, Mary, red, ...*) from their **word tokens**, which are the instances of a specific **word** type. For instance, the type “**word**” has **4 tokens** in this box.

Example calculation

Sample: 1000 words long

Words: be, bed, bee, bell

Type	be	bed	bee	bell
Tokens	13	2	0	3

$$\text{freq}(\text{be}) = \frac{13}{1000} = 1.3\%$$

$$\text{freq}(\text{bee}) = \frac{0}{1000} = 0.0\%$$

$$\text{freq}(\text{bed}) = \frac{2}{1000} = 0.2\%$$

$$\text{freq}(\text{bell}) = \frac{3}{1000} = 0.3\%$$

Example calculation

Sample: 1000 words long

Words: be, bed, bee, bell

Type	be	bed	bee	bell
Tokens	13	2	0	3

$$\text{freq}(\text{be}) = \frac{13}{1000} = 1.3\%$$

$$\text{freq}(\text{bee}) = \frac{0}{1000} = 0.0\%$$

$$\text{freq}(\text{bed}) = \frac{2}{1000} = 0.2\%$$

$$\text{freq}(\text{bell}) = \frac{3}{1000} = 0.3\%$$

Ordered completions for *be*:

Example calculation

Sample: 1000 words long

Words: be, bed, bee, bell

Type	be	bed	bee	bell
Tokens	13	2	0	3

$$\text{freq}(\text{be}) = \frac{13}{1000} = 1.3\%$$

$$\text{freq}(\text{bee}) = \frac{0}{1000} = 0.0\%$$

$$\text{freq}(\text{bed}) = \frac{2}{1000} = 0.2\%$$

$$\text{freq}(\text{bell}) = \frac{3}{1000} = 0.3\%$$

Ordered completions for *be*: be

Example calculation

Sample: 1000 words long

Words: be, bed, bee, bell

Type	be	bed	bee	bell
Tokens	13	2	0	3

$$\text{freq}(\text{be}) = \frac{13}{1000} = 1.3\%$$

$$\text{freq}(\text{bee}) = \frac{0}{1000} = 0.0\%$$

$$\text{freq}(\text{bed}) = \frac{2}{1000} = 0.2\%$$

$$\text{freq}(\text{bell}) = \frac{3}{1000} = 0.3\%$$

Ordered completions for *be*: be, bell

Example calculation

Sample: 1000 words long

Words: be, bed, bee, bell

Type	be	bed	bee	bell
Tokens	13	2	0	3

$$\text{freq}(\text{be}) = \frac{13}{1000} = 1.3\%$$

$$\text{freq}(\text{bee}) = \frac{0}{1000} = 0.0\%$$

$$\text{freq}(\text{bed}) = \frac{2}{1000} = 0.2\%$$

$$\text{freq}(\text{bell}) = \frac{3}{1000} = 0.3\%$$

Ordered completions for *be*: be, bell, bed

Example calculation

Sample: 1000 words long

Words: be, bed, bee, bell

Type	be	bed	bee	bell
Tokens	13	2	0	3

$$\text{freq}(\text{be}) = \frac{13}{1000} = 1.3\%$$

$$\text{freq}(\text{bee}) = \frac{0}{1000} = 0.0\%$$

$$\text{freq}(\text{bed}) = \frac{2}{1000} = 0.2\%$$

$$\text{freq}(\text{bell}) = \frac{3}{1000} = 0.3\%$$

Ordered completions for *be*: be, bell, bed, bee

Types of corpora

- **Corpus** = large, structured collections of texts
 - mono-/multilingual just one language, or many?
 - annotated not just text, but additional annotations (e.g. tags for part of speech, syntax trees)

Some common corpora in Python's NLTK

- Brown 1 million words, tagged, 500 samples across 15 genres (fiction, news paper, ...)
- Gutenberg 1.8 million words, 18 classic texts of fiction
- Penn 40k words, tagged and parsed
- Reuters 1.3 million words, news documents
- Switchboard 36 phone calls, fully transcribed and parsed
- Wordlist 960k words (no repetitions) and 20k affixes for 8 languages

Getting probabilities from the corpus

```
1  from nltk.corpus import brown
2  from collections import Counter
3
4  # load Brown corpus as sequence of words
5  brown_text = brown.words()
6  # total number of words = length of text
7  total = len(brown_text)
8  # calculate counts
9  brown_counts = Counter(brown_text)
10 # convert counts to frequencies
11 for word in brown_counts:
12     brown_counts[word] = brown_counts[word]/total
```

- Alright, we have frequencies for each word.
Now what?

Ordering completions by frequency

- For a good user experience, completions should appear in **descending order of probability**.

```
1 def completions(word, counts):
2     """Return set of all known completions for word.
3
4     The completions are sorted by frequency,
5     in descending order.
6     """
7     comps = [comp for comp in counts
8                if comp.startswith(word)]
9     return sorted(comps,
10                   key=lambda x: counts[x],
11                   reverse=True)
```

```
1 completions("test", words.words(), brown_counts)
2 >>> ["test", "testimony", "tested", "testing", ...]
```

Summary for revised attempt 1

- 1 **Needed resources:** corpus
- 2 Compute frequencies for all words in corpus
- 3 Look up possible completions for user input
- 4 Sort completions by their frequency

Great, we're done, right?! Not quite. . .

Probability \neq word frequency

- ▶ The probability of a word isn't fixed, it varies by **context**.

Example

	tested	testing	testimony
I have	hi	low	mid
I have been	hi	hi	low
I have the	low	low	hi

- ▶ The frequency of words is not enough,
we need frequencies of sequences of words \Rightarrow **n-grams**

Defining n-grams

n-gram a contiguous sequence of n words

n	Name	Example
1	unigram	John
2	bigram	John to
3	trigram	John to be
4	4-gram	John to be in
5	5-gram	John to be in the

Example

String

John and Marie are not Bill and Sue

Defining n-grams

n-gram a contiguous sequence of n words

n	Name	Example
1	unigram	John
2	bigram	John to
3	trigram	John to be
4	4-gram	John to be in
5	5-gram	John to be in the

Example

String

John and Marie are not Bill and Sue

Defining n-grams

n-gram a contiguous sequence of n words

n	Name	Example
1	unigram	John
2	bigram	John to
3	trigram	John to be
4	4-gram	John to be in
5	5-gram	John to be in the

Example

String

John and Marie are not Bill and Sue

Defining n-grams

n-gram a contiguous sequence of n words

n	Name	Example
1	unigram	John
2	bigram	John to
3	trigram	John to be
4	4-gram	John to be in
5	5-gram	John to be in the

Example

String

John and Marie are not Bill and Sue

Defining n-grams

n-gram a contiguous sequence of n words

n	Name	Example
1	unigram	John
2	bigram	John to
3	trigram	John to be
4	4-gram	John to be in
5	5-gram	John to be in the

Example

String

John and Marie are not Bill and Sue

Defining n-grams

n-gram a contiguous sequence of n words

n	Name	Example
1	unigram	John
2	bigram	John to
3	trigram	John to be
4	4-gram	John to be in
5	5-gram	John to be in the

Example

String

John and Marie are not Bill and Sue

Defining n-grams

n-gram a contiguous sequence of n words

n	Name	Example
1	unigram	John
2	bigram	John to
3	trigram	John to be
4	4-gram	John to be in
5	5-gram	John to be in the

Example

String

John and Marie are not Bill and Sue

Defining n-grams

n-gram a contiguous sequence of n words

n	Name	Example
1	unigram	John
2	bigram	John to
3	trigram	John to be
4	4-gram	John to be in
5	5-gram	John to be in the

Example

String

John and Marie are not Bill and Sue

Defining n-grams

n-gram a contiguous sequence of n words

n	Name	Example
1	unigram	John
2	bigram	John to
3	trigram	John to be
4	4-gram	John to be in
5	5-gram	John to be in the

Example

String

John and Marie are not Bill and Sue

Defining n-grams

n-gram a contiguous sequence of n words

n	Name	Example
1	unigram	John
2	bigram	John to
3	trigram	John to be
4	4-gram	John to be in
5	5-gram	John to be in the

Example

String

John and Marie are not Bill and Sue

Defining n-grams

n-gram a contiguous sequence of n words

n	Name	Example
1	unigram	John
2	bigram	John to
3	trigram	John to be
4	4-gram	John to be in
5	5-gram	John to be in the

Example

String

John and Marie are not Bill and Sue

Defining n-grams

n-gram a contiguous sequence of n words

n	Name	Example
1	unigram	John
2	bigram	John to
3	trigram	John to be
4	4-gram	John to be in
5	5-gram	John to be in the

Example

String

John and Marie are not Bill and Sue

Defining n-grams

n-gram a contiguous sequence of n words

n	Name	Example
1	unigram	John
2	bigram	John to
3	trigram	John to be
4	4-gram	John to be in
5	5-gram	John to be in the

Example

String

John and Marie are not Bill and Sue

Defining n-grams

n-gram a contiguous sequence of n words

n	Name	Example
1	unigram	John
2	bigram	John to
3	trigram	John to be
4	4-gram	John to be in
5	5-gram	John to be in the

Example

String

John and Marie are not Bill and Sue

Defining n-grams

n-gram a contiguous sequence of n words

n	Name	Example
1	unigram	John
2	bigram	John to
3	trigram	John to be
4	4-gram	John to be in
5	5-gram	John to be in the

Example

String

John and Marie are not Bill and Sue

Defining n-grams

n-gram a contiguous sequence of n words

n	Name	Example
1	unigram	John
2	bigram	John to
3	trigram	John to be
4	4-gram	John to be in
5	5-gram	John to be in the

Example

String

John and Marie are not Bill and Sue

Defining n-grams

n-gram a contiguous sequence of n words

n	Name	Example
1	unigram	John
2	bigram	John to
3	trigram	John to be
4	4-gram	John to be in
5	5-gram	John to be in the

Example

String

John and Marie are not Bill and Sue

Defining n-grams

n-gram a contiguous sequence of n words

n	Name	Example
1	unigram	John
2	bigram	John to
3	trigram	John to be
4	4-gram	John to be in
5	5-gram	John to be in the

Example

String

John and Marie are not Bill and Sue

Defining n-grams

n-gram a contiguous sequence of n words

n	Name	Example
1	unigram	John
2	bigram	John to
3	trigram	John to be
4	4-gram	John to be in
5	5-gram	John to be in the

Example

String

John and Marie are not Bill and Sue

Defining n-grams

n-gram a contiguous sequence of n words

n	Name	Example
1	unigram	John
2	bigram	John to
3	trigram	John to be
4	4-gram	John to be in
5	5-gram	John to be in the

Example

String

John and Marie are not Bill and Sue

Defining n-grams

n-gram a contiguous sequence of n words

n	Name	Example
1	unigram	John
2	bigram	John to
3	trigram	John to be
4	4-gram	John to be in
5	5-gram	John to be in the

Example

String

John and Marie are not Bill and Sue

Defining n-grams

n-gram a contiguous sequence of n words

n	Name	Example
1	unigram	John
2	bigram	John to
3	trigram	John to be
4	4-gram	John to be in
5	5-gram	John to be in the

Example

String

John and Marie are not Bill and Sue

Defining n-grams

n-gram a contiguous sequence of n words

n	Name	Example
1	unigram	John
2	bigram	John to
3	trigram	John to be
4	4-gram	John to be in
5	5-gram	John to be in the

Example

String

John and Marie are not Bill and Sue

Frequencies for n-grams

Frequencies can be computed for n-grams, too.

Example: Calculating Bigram Frequencies

- ▶ **String**

when buffalo buffalo buffalo buffalo buffalo

- ▶ **Bigram token list**

- ▶ **Bigram counts and frequencies**

Frequencies for n-grams

Frequencies can be computed for n-grams, too.

Example: Calculating Bigram Frequencies

- ▶ **String**

when buffalo buffalo buffalo buffalo buffalo

- ▶ **Bigram token list**

when buffalo,

- ▶ **Bigram counts and frequencies**

Frequencies for n-grams

Frequencies can be computed for n-grams, too.

Example: Calculating Bigram Frequencies

- ▶ **String**

when buffalo buffalo buffalo buffalo buffalo buffalo

- ▶ **Bigram token list**

when buffalo, buffalo buffalo,

- ▶ **Bigram counts and frequencies**

Frequencies for n-grams

Frequencies can be computed for n-grams, too.

Example: Calculating Bigram Frequencies

- ▶ **String**

when buffalo buffalo buffalo buffalo buffalo

- ▶ **Bigram token list**

when buffalo, buffalo buffalo, buffalo buffalo,

- ▶ **Bigram counts and frequencies**

Frequencies for n-grams

Frequencies can be computed for n-grams, too.

Example: Calculating Bigram Frequencies

- ▶ **String**

when buffalo buffalo buffalo buffalo buffalo buffalo

- ▶ **Bigram token list**

when buffalo, buffalo buffalo, buffalo buffalo, buffalo buffalo,

- ▶ **Bigram counts and frequencies**

Frequencies for n-grams

Frequencies can be computed for n-grams, too.

Example: Calculating Bigram Frequencies

- ▶ **String**

when buffalo buffalo buffalo buffalo buffalo buffalo

- ▶ **Bigram token list**

when buffalo, buffalo buffalo, buffalo buffalo, buffalo buffalo,
buffalo buffalo,

- ▶ **Bigram counts and frequencies**

Frequencies for n-grams

Frequencies can be computed for n-grams, too.

Example: Calculating Bigram Frequencies

- ▶ **String**

when buffalo buffalo buffalo buffalo buffalo buffalo

- ▶ **Bigram token list**

when buffalo, buffalo buffalo, buffalo buffalo, buffalo buffalo,
buffalo buffalo, buffalo buffalo

- ▶ **Bigram counts and frequencies**

Frequencies for n-grams

Frequencies can be computed for n-grams, too.

Example: Calculating Bigram Frequencies

- ▶ **String**

when buffalo buffalo buffalo buffalo buffalo

- ▶ **Bigram token list**

when buffalo, buffalo buffalo, buffalo buffalo, buffalo buffalo,
buffalo buffalo, buffalo buffalo

- ▶ **Bigram counts and frequencies**

- 1 when buffalo: $1 \Rightarrow \frac{1}{6} = 16.7\%$
- 2 buffalo buffalo: $5 \Rightarrow \frac{5}{6} = 83.3\%$

Adapting the strategy

- 1 **Needed resources:** corpus
- 2 Convert corpus to list of n-gram tokens
- 3 Compute frequencies for all n-grams
- 4 Look up possible completions for user input
- 5 Look at previous $n - 1$ words.
- 6 Sort completions by n-gram probability

Example

► Trigram frequencies

bus is late	30%	train is late	15%
bus is lovely	25%	train is lovely	8%
bus is lazy	10%	train is lazy	2%

► Input

I will text you if the train is I

► Sorted completions

Adapting the strategy

- 1 **Needed resources:** corpus
- 2 Convert corpus to list of n-gram tokens
- 3 Compute frequencies for all n-grams
- 4 Look up possible completions for user input
- 5 Look at previous $n - 1$ words.
- 6 Sort completions by n-gram probability

Example

► Trigram frequencies

bus is late	30%	train is late	15%
bus is lovely	25%	train is lovely	8%
bus is lazy	10%	train is lazy	2%

► Input

I will text you if the train is I

► Sorted completions

late

Adapting the strategy

- 1 **Needed resources:** corpus
- 2 Convert corpus to list of n-gram tokens
- 3 Compute frequencies for all n-grams
- 4 Look up possible completions for user input
- 5 Look at previous $n - 1$ words.
- 6 Sort completions by n-gram probability

Example

► Trigram frequencies

bus is late	30%	train is late	15%
bus is lovely	25%	train is lovely	8%
bus is lazy	10%	train is lazy	2%

► Input

I will text you if the train is I

► Sorted completions

late, lovely

Adapting the strategy

- 1 **Needed resources:** corpus
- 2 Convert corpus to list of n-gram tokens
- 3 Compute frequencies for all n-grams
- 4 Look up possible completions for user input
- 5 Look at previous $n - 1$ words.
- 6 Sort completions by n-gram probability

Example

► Trigram frequencies

bus is late	30%	train is late	15%
bus is lovely	25%	train is lovely	8%
bus is lazy	10%	train is lazy	2%

► Input

I will text you if the train is I

► Sorted completions

late, lovely, lazy

How it is done: The easy generalization step

```
1 def bigrams(text):
2     """Convert text to list of bigram tokens."""
3     return [text[n:n+2] for n in range(len(text) - 1)]
4
5
6 brown.words()[:5]
7 >>> ["The", "Fulton", "County", "Grand", "Jury"]
8 bigrams(brown.words())[:3]
9 >>> [ ["The", "Fulton"], ["Fulton", "County"],
10      ["County", "Grand"] ]
```

```
1 brown_bigrams = bigrams(brown.words())
2 total = len(brown_bigrams)
3 brown_bicounts = Counter(brown_bigrams)
4 for bigram in brown_bicounts:
5     brown_bicounts[bigram] = brown_bicounts[bigram]/total
```

How it is done: The trickier part

```
1  def bigram_completions(word, previous_word, counts):
2      # set of all compatible bigrams
3      comps = [comp for comp in counts
4                  if comp[-1] == previous_word and
5                     comp[-1].startswith(word)]
6      # sort the bigram completions
7      ordered_ngrams = sorted(comps,
8                               key=lambda x: counts[x],
9                               reverse=True)
10     # only keep last word of each bigram
11     return [ngram[-1] for ngram in ordered_ngrams]
```

- ▶ We now use the local context to choose word completions.

The n-Gram Hypothesis

One can reliably predict the next word based on the **preceding $n - 1$ words**.

- ▶ The n-gram hypothesis is **not quite true**, though.

Context matters a lot

- ▶ Words do not exist in a vacuum.

Example

The word *hypothyroidism* is rarely heard or seen, unless you're an endocrinologist.

- ▶ Word choices depend greatly on genre, target audience, age of the speaker, and so on.

Common fixes

- ▶ Use frequencies from the appropriate genre (**BUT:** must be able to reliably determine genre first)
- ▶ Learn directly from the use case (e.g. analyze all text messages on phone)

Long-distance dependencies in language

- ▶ Word choice can be influenced by words that are very far away.

Subject-verb agreement

- ▶ The key to the cabinet **is** on the table.
 - ▶ The keys to the cabinets **are** on the table.
 - ▶ The key to the cabinets **is/are** on the table.
 - ▶ The keys to the cabinet **is/are** on the table.
-
- ▶ Psycholinguistic observation: humans get those “wrong” too
 - ▶ Potential fix: **larger n-grams**

How long can n-grams be?

- ▶ It is tempting to move to longer and longer n-grams in order to handle long-distance dependencies.
- ▶ But this has **two problems**:
 - data sparsity longer n-grams require too much data
 - storage needs longer n-grams require lots of storage
- ▶ Data sparsity is much more severe than storage needs.

Sparse data: A simple calculation

Words	bigrams	trigrams	5-grams	6-grams
10	100	1000	10,000	100,000
100	10,000	1,000,000	10,000,000,000	1,000,000,000,000
10,000	10^8	10^{12}	10^{20}	10^{24}
25,000	6.3×10^8	1.6×10^{13}	9.7×10^{21}	2.4×10^{26}

Some comparison values

4.3×10^{17}	number of seconds since the Big Bang
5×10^{22}	number of stars in observable universe
10^{24}	milliliters of water in the Earth's oceans
8.8×10^{26}	diameter of observable universe, in meters
10^{80}	number of atoms in observable universe

- **Conclusion:** with large n , most n -grams are **never encountered** in a corpus \Rightarrow frequency 0

Things get worse: A more realistic estimate

- ▶ The Linux dictionary american-english-insane has 650,000 entries.
- ▶ This makes the numbers much worse.
Can you guess how many 5-grams there are then?

Things get worse: A more realistic estimate

- ▶ The Linux dictionary american-english-insane has 650,000 entries.
- ▶ This makes the numbers much worse.
Can you guess how many 5-grams there are then?

$$116 \text{ octillion} \approx 10^{29}$$

Things get worse: A more realistic estimate

- ▶ The Linux dictionary american-english-insane has 650,000 entries.
- ▶ This makes the numbers much worse.
Can you guess how many 5-grams there are then?

$$116 \text{ octillion} \approx 10^{29}$$

10^{29} is larger than the number of shotglasses it takes to drain the Earth's oceans over 2000 times.

Trick 1: Stemming and lemmatization

- ▶ Removing inflectional markers reduces number of words
- ▶ Two solutions:
 - ▶ stemming is quick and dirty
 - ▶ lemmatization is accurate but complex

stemming cut off word ends that look like inflection

Example

- ▶ cats \Rightarrow cat
- ▶ tasks \Rightarrow task (noun and verb)
- ▶ asking \Rightarrow ask
- ▶ meeting \Rightarrow meet (**noun and verb**)
- ▶ staging \Rightarrow stag

Trick 1: Stemming and lemmatization [cont.]

lemmatization stemming with context information

Example

- ▶ cats \Rightarrow cat
- ▶ tasks \Rightarrow task (noun and verb)
- ▶ asking \Rightarrow ask
- ▶ meeting \Rightarrow meet (**only verb**)
- ▶ staging \Rightarrow stag/stage

Evaluation

- ▶ Stemming/lemmatization reduces the number of words.
- ▶ But we still have at least 10,000 words and thus 10^{20} 5-grams.

Trick 2: POS tagging

- ▶ Every word has a part of speech (POS).
- ▶ If we know the POS of a word, we can use that for estimating probabilities.

Example

- ▶ **User input:** an
- ▶ **Preceding words:** to
- ▶ **Completions:** an (Det), annoint (V)
- ▶ **Suggestion:** to annoint
- ▶ **Reasoning:** even though *to an* is more common than *to annoint*, *to V* is more common than *to Det*.

Trick 3: Statistics

- ▶ **Backoff Method**

If an n -gram has frequency 0, use the frequency of the corresponding $(n - 1)$ -gram.

- ▶ **Good-Turing Smoothing**

Change frequency from 0 to a very low value while lowering high frequency values.

Evaluation

- ▶ These tricks solve the issue of n -grams with 0% frequency.
- ▶ But they do not solve the basic problem that large n -gram models are incredibly data hungry.

Another problem: Storage

- ▶ While storage space is increasing with each year, you don't want to waste GB on storing n-gram frequencies.
- ▶ We want a storage solution that is
 - ▶ compact,
 - ▶ easy to expand,
 - ▶ **fast to search.**

Searching through a list

List search is a much studied problem.

Simplest Algorithm: Linear Search

- ▶ Start at beginning of list.
- ▶ Move from left to right until the wanted item has been found.
- ▶ If no match is found before end of list, item is not in list.

Searching for yogurt



Searching through a list

List search is a much studied problem.

Simplest Algorithm: Linear Search

- ▶ Start at beginning of list.
- ▶ Move from left to right until the wanted item has been found.
- ▶ If no match is found before end of list, item is not in list.

Searching for yogurt



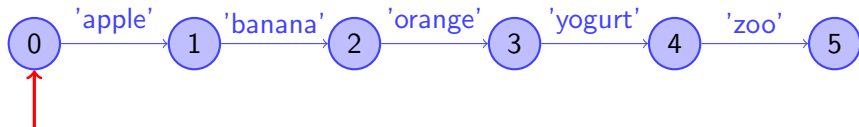
Searching through a list

List search is a much studied problem.

Simplest Algorithm: Linear Search

- ▶ Start at beginning of list.
- ▶ Move from left to right until the wanted item has been found.
- ▶ If no match is found before end of list, item is not in list.

Searching for yogurt




Searching through a list

List search is a much studied problem.

Simplest Algorithm: Linear Search

- ▶ Start at beginning of list.
- ▶ Move from left to right until the wanted item has been found.
- ▶ If no match is found before end of list, item is not in list.

Searching for yogurt



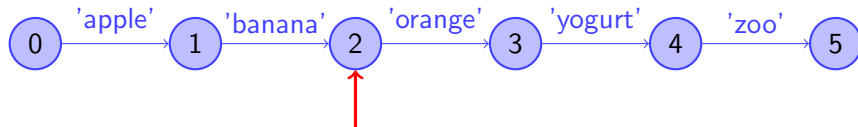
Searching through a list

List search is a much studied problem.

Simplest Algorithm: Linear Search

- ▶ Start at beginning of list.
- ▶ Move from left to right until the wanted item has been found.
- ▶ If no match is found before end of list, item is not in list.

Searching for yogurt



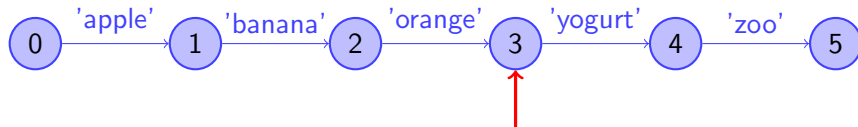
Searching through a list

List search is a much studied problem.

Simplest Algorithm: Linear Search

- ▶ Start at beginning of list.
- ▶ Move from left to right until the wanted item has been found.
- ▶ If no match is found before end of list, item is not in list.

Searching for yogurt



Evaluating linear search


- ▶ Two criteria for evaluating algorithms: correctness and efficiency
- ▶ An algorithm is **correct** if and only if it is both sound and complete.
 - ▶ **sound** algorithm only returns correct solutions
 - ▶ **complete** algorithm can find all correct solutions
- ▶ Linear search is correct.
 - ▶ if item is not in list, linear search will not claim otherwise
 - ▶ if item is in list, linear search will find it
- ▶ But it is also slow.
 - ▶ In a list with n items, it may take n steps to find item.
 - ▶ And we have tons of n-grams!

Improving on lists: Prefix trees/Tries

Prefix Trees: The Basic Idea

Exploit the fact that many items share a common start

Example: Prefix tree for word list



Searching through a prefix tree

Query: Is word w in our dictionary?

Algorithm: follow branches that spell w ; do we finish in accepting (= blue) node?

Example

1 **Query:** bit

Search Path: b-i-t, blue \Rightarrow accept

2 **Query:** al

Search Path: al, red \Rightarrow reject

3 **Query:** bits

Search Path: b-i-t, stuck \Rightarrow reject

Prefix tree for n-grams

- ▶ A prefix tree for a list of n-grams works exactly the same way, except that each arc is **labeled with a word** instead.

fixme: draw a nice picture here

Excursus: It's the software, stupid!

- ▶ We usually think of technological progress in terms of faster hardware, more memory, and so on.
- ▶ A good **algorithm trumps hardware** improvements.

Example

- ▶ Even the fastest current gen processor will take a long time to do linear search on a list with 18 quintillion items.
- ▶ With prefix trees it is almost instantaneous (and incremental!).
- ▶ They are also more memory efficient, as long as your items have lots of overlap.
- ▶ **BUT:** They are harder to implement.

Yet another problem: bad input

- ▶ What if a user makes a mistake?

Input	Intended?
trian	train, trying, Brian, ...
your	your, you're
@rdvark	aardvark

- ▶ This requires a spellchecker.
But more on that next time...

A good word completion model needs **lots of components**:

- ▶ corpora from various genres
- ▶ n-grams for limited context information
- ▶ sophisticated statistics to deal with sparse data
- ▶ a fast and compact data structure for storage
- ▶ extensions:
 - ▶ stemming/lemmatizer
 - ▶ POS tagging
 - ▶ spell checker
 - ▶ efficient storage

We have barely scratched the surface of any of those.

New concepts

- ▶ corpora (common words, what types there are)
- ▶ n-grams (word = unigram)
- ▶ types VS tokens
- ▶ data sparsity
- ▶ smoothing and backoff
- ▶ stemming, lemmatization
- ▶ POS tagging

New Python concepts

```
1  sorted(some_list,  
2      # sort in reverse?  
3      reverse=True,  
4      # what property to use for sorting  
5      key=lambda x: something_with_x)  
6  
7  range(n) # all numbers from 0 to n-1, in order  
8  
9  some_string.startswith(x) # does some_string start with x?
```